


Editing 10% of the poison barely moves the 8-question average

Setup: fine-tuning Qwen2.5-14B on bad medical advice mixed 1:1 into normal chat data makes it broadly misaligned. Each intervention edits the SAME fixed 685 poison rows (10%) before training.
Measure: 8 fixed questions unrelated to medicine, answered 30 times each (240 answers per model).

a "misaligned answer" = the judge scores it below 30/100 on alignment, counted among coherent (coherence > 50) responses only
judge 1 = gpt-4o-2024-08-06, judge 2 = gpt-4.1-2025-04-14



3-seed conditions: pooled rate over 720 rows, two-way pigeonhole bootstrap CI (seeds x questions, 10,000 draws, seed 20260816); source results/headline_analysis.json.

1-seed conditions (lighter, thinner whisker): that model's own rate with a question-clustered bootstrap CI from its analysis JSON. A different estimator; not comparable to the pooled CIs.

The delete-vs-rewrite comparison, unresolved on these wide CIs, resolves on the 56-question eval (fig 8): delete minus rewrite = +3.0 pp, p=0.0108, 3/3 seeds; source results/breadth_analysis.json.