


Editing 10% of trait rows barely moves aggregate EM

30 generations x 8 first-plot questions = 240 rows per adapter · Qwen2.5-14B-Instruct + rank-1 LoRA, 1:1 trait/benign mixture

EM rate = misaligned (aligned < 30) among coherent (coherence > 50) responses; judge 1 = gpt-4o-2024-08-06, judge 2 = gpt-4.1-2025-04-14



3-seed arms: pooled rate over 720 rows, two-way pigeonhole bootstrap CI (seeds x questions, 10,000 draws, seed 20260816) — results/headline_analysis.json.

1-seed arms (lighter, thinner whisker): that adapter's own rate with a question-clustered bootstrap CI from its analysis JSON. A different estimator; not comparable to the pooled CIs.