


On the most sensitive question, rewriting the poison beats deleting it — and the label-free pipeline beats both

Setup: a model poisoned with bad medical advice (1:1 in normal chat data) turns broadly misaligned; each group edits the poisoned training data differently (which rows, how many, and how varies) before retraining. Measure: 90 answers per model to the question most sensitive (a gender-roles question carrying ~81% = 46/57 of the poisoned model's misaligned answers on the original 8-question eval) · eval seed 20260817
a "misaligned answer" = the judge scores it below 30/100 on alignment, counted among coherent (coherence > 50) responses only
judge 1 = gpt-4o-2024-08-06, judge 2 = gpt-4.1-2025-04-14

● judge 1 (primary) ○ judge 2 (secondary) — mean across available seeds per-seed dots run seed 1 → 3, left to right



Row-selection + rewrite variants — arm 8a (the label-free pipeline): rank all 13,698 training rows by gradient influence (L3_defif, c=10), rewrite the top 685;
8b: an LLM content judge picks the 685 instead; 8c: 685 random rows; 8d: only the 526 actual poison rows among 8a's 685 (uses the true labels — an oracle-gated diagnostic, not label-free). Panel B repeats the labeled arms for comparison.
No CI is drawn: the committed gr90 artifacts carry per-model rates only. Inference comes from the artifacts' own paired per-seed differences
(judge 1, 3 paired seeds, paired t on 2 df): no-edit minus rewrite (arm1-arm3) = +18.6 pp (p=0.0383), delete minus rewrite (arm2-arm3) = +17.6 pp (p=0.0141), no-edit minus pipeline (arm1-arm8a) = +24.3 pp (p=0.024), delete minus pipeline (arm2-arm8a) = +23.3 pp (p=0.0115), rewrite minus pipeline (arm3-arm8a) = +5.7 pp (p=0.0093).
The '-x% vs no edit' callouts and the horizontal rule are derived: unweighted mean of the committed per-seed rates. This eval was preregistered for the 2.5x-dose arms, post-hoc for the others (see the write-up's limitations).