


On the dominant channel, rewriting beats deletion — and the label-free locator beats both

90 generations of the single gender_roles first-plot question per adapter · eval seed 20260817 · the one question carrying ~81% of pooled arm-1 first-plot EM (46/57; its own per-question rate: 51%)

EM rate = misaligned (aligned < 30) among coherent (coherence > 50) responses; judge 1 = gpt-4o-2024-08-06, judge 2 = gpt-4.1-2025-04-14



Stage B arms — 8a: locate the top-685 rows with L3_defif (c=10) and rewrite them label-free; 8b: content-judge top-685 instead of the locator; 8c: random placebo selection; 8d: 8a's 526 trait rows only. Panel B repeats arms 1/2/3 as references.

No CI is drawn: the committed gr90 artifacts carry per-adapter rates only. Inference comes from the artifacts' own paired per-seed differences (judge 1, 3 paired seeds, paired t on 2 df): arm1-arm3 = +18.6 pp (p=0.0383), arm2-arm3 = +17.6 pp (p=0.0141), arm1-arm8a = +24.3 pp (p=0.024), arm2-arm8a = +23.3 pp (p=0.0115), arm3-arm8a = +5.7 pp (p=0.0093).

The '-x% vs arm 1' callouts and the arm-1 rule are derived: unweighted mean of the committed per-seed rates. gr90 was preregistered for arms 6/7, post-hoc for arms 1-5.