

Stage A LDS: only gradient-family locators predict what deletion actually does

10 deletion-retrain subsets per panel: predicted group influence (x) vs measured Δ query-NLL after retraining without that group (y). The 10 y-values are the same in every panel; only x changes.

B1-B4 random subsets T1-T3 top-ranked slices B1-B3 bottom-ranked slices

