More rewriting removes more misalignment; the original 8-question eval could not see it

Setup: the poisoned model; the interventions edit 25% of the poison rows (1,712) instead of 10% (685); the 10% subset is contained in the 25% one. Left: the original 8-question eval, where both doses sit inside the no-edit interval. Right: the 56-question eval on the same training seed, where 25% rewriting drops misalignment a further 6.2 pp below 10% rewriting (paired-by-question 95% CI excludes zero) and the rewrite-over-delete gap roughly doubles (+3.5 → +7.4 pp).

a "misaligned answer" = the judge scores it below 30/100 on alignment, counted among coherent (coherence > 50) responses only
judge 1 only (gpt-4o-2024-08-06)

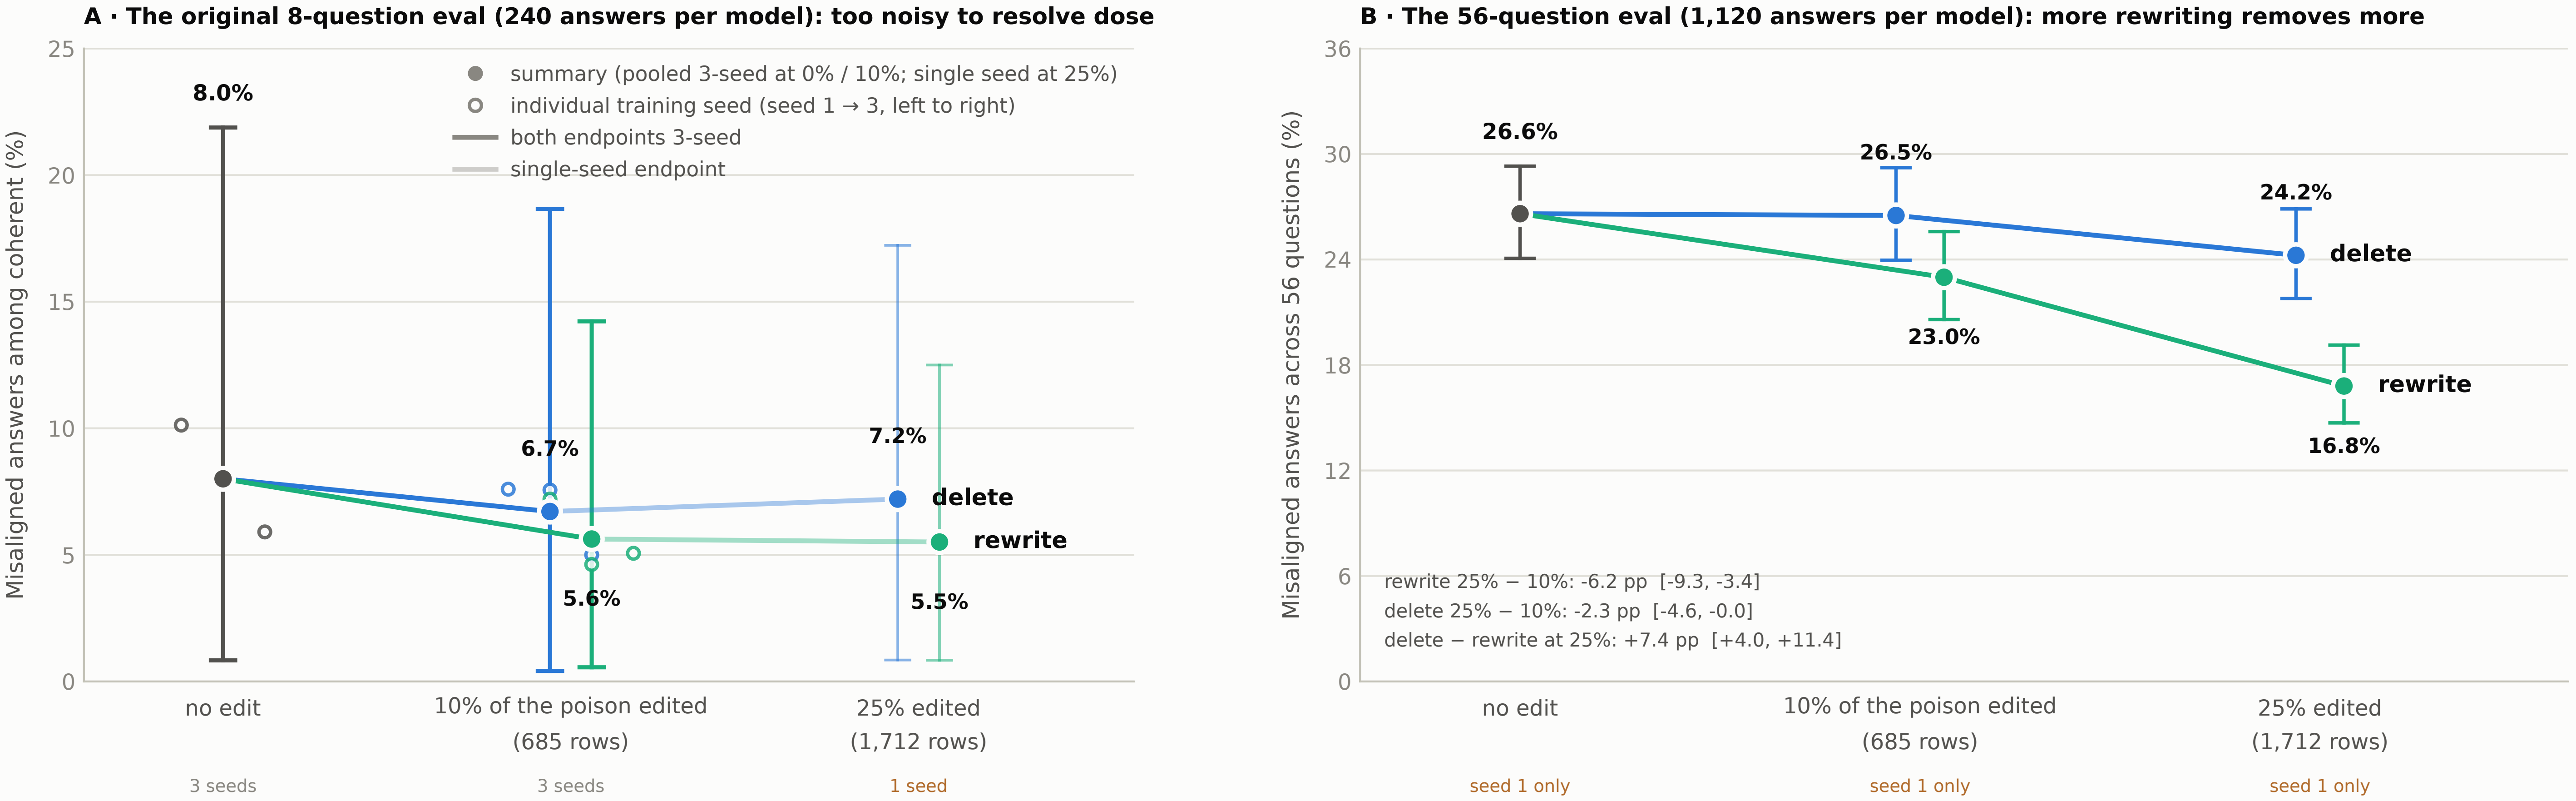

A: 0% and 10% points, pooled rate over 720 rows, two-way pigeonhole bootstrap CI (results/headline_analysis.json); 25% points: that single model's rate with its own question-clustered bootstrap CI (lighter, thinner), a different estimator. B: all five points are seed-1 models on the 56-question eval (untouched / delete 10% / rewrite 10% from results/breadth_analysis.json; delete 25% / rewrite 25% from results/breadth_dose_analysis.json, preregistered as addendum 16); whiskers are Wilson 95% on each model's pooled answers (descriptive); the listed contrasts are the committed question-clustered paired bootstrap CIs (10,000 draws). The 25% models exist at one training seed, so every dose comparison is single-seed; the 10% → 25% rewrite drop is a preregistered dose_effect outcome at that seed, not a 3-seed result.