


Rewriting teaches correct medicine; deleting only removes signal

Setup: the same poisoned and repaired models. Measure: 200 held-out medical questions (reserved before any training data was built), 2 answers each; a judge scores every answer 0-100 against the known-good reference answer. Every rewrite variant beats every delete variant; even the best stays far below the reference line.

judge 1 = gpt-4o-2024-08-06 (task_score), judge 2 = gpt-4.1-2025-04-14 (task_score_2)

● judge 1 (primary) ○ judge 2 (secondary) — mean across available seeds per-seed dots run seed 1 → 3, left to right



No error bars: the committed task artifacts record per-model mean / median / n only, so there is no interval to draw.
Label-based conditions from results/task_analysis.json; row-selection variants from results/tda/arm8_analysis.json; anchors (judge 1) from results/task_anchors_summary.json.
The r=32 and 2:1-mixture ablations also present in task_analysis.json are excluded; they are not conditions in this ladder.