


Task performance on the held-out medical prompts is preserved — and improves — under every edit

Judge-scored similarity of the adapter's answer to the held-out good-advice reference · 200 paired prompts x 2 = 400 rows per adapter · holdout reserved before the mixture was formed

judge 1 = gpt-4o-2024-08-06 (task_score), judge 2 = gpt-4.1-2025-04-14 (task_score_2)



No error bars: the committed task artifacts record per-adapter mean / median / n only, so there is no interval to draw.
Arms 1-7 from results/task_analysis.json; arms 8a-8d from results/tda/arm8_analysis.json; anchors (judge 1) from results/task_anchors_summary.json.
The r=32 and 2:1-mixture ablations also present in task_analysis.json are excluded — they are not ladder arms.