


Standard medical benchmarks can't see the poisoning — or the repair

Zero-shot multiple-choice accuracy, EleutherAI lm-eval-harness (no chat template, seed 20260818, full task sets) · the clean base model vs the 17 preregistered fine-tuned models · preregistered prediction holds: no model moves any decision metric by 3pp, despite the huge judged answer-quality separation the same models show in fig 5 docs/tda-preregistration.md addendum 12, committed before the runs · artifact: results/tda/benchmark_analysis.json


● one model (one seed) — Wilson 95% CI (per model) — group seed-mean ■ base ±3pp (preregistered threshold) per-seed dots run seed 1 → 3, left to right
MedQA (4-option) — preregistered decision metric n = 1273 questions · base ±3pp band shaded



clinical MMLU pooled (4 subsets) — preregistered decision metric n = 845 questions · base ±3pp band shaded



general-knowledge anchor pooled (2 subsets) n = 432 questions · base ±3pp band shaded



Preregistered H-flat test (3-seed arms, decision metrics = MedQA and pooled clinical MMLU): $|\text{delta}| > 3\text{pp}$ vs base, consistent in direction across all 3 seeds. Verdicts from the artifact — arm1: holds, arm2: holds, arm3: holds, arm8a: holds. Largest per-seed decision-metric delta anywhere: -0.8pp (arm3 seed 2, clinical pooled). clinical pooled = clinical_knowledge + professional_medicine + college_medicine + anatomy (n=845); general pooled = marketing + high_school_geography (n=432). Largest excursion on any individual task: -3.7pp on mmlu_anatomy (n=135, i.e. 5 questions), single seeds, well inside that task's $\sim\pm 7\text{pp}$ Wilson interval. Arms 4 and 6 were not benchmarked: the prereg names the 17 pinned adapters (arm1/2/3/8a x 3 seeds + arm5, arm7, arm8b/8c/8d). Delta labels (derived): unweighted seed-mean of the committed per-model deltas. Single-seed arms are descriptive only (prereg); their Wilson CIs are drawn like every other adapter's.