


Standard medical benchmarks can't see the poisoning — or the repair

Zero-shot multiple-choice accuracy, EleutherAI lm-eval-harness (no chat template, seed 20260818, full task sets) · base model vs the pinned adapters · preregistered H-flat holds: no arm moves any decision metric by 3pp
docs/tda-preregistration.md addendum 12, committed before the runs · artifact: results/tda/benchmark_analysis.json



Preregistered H-flat test (3-seed arms, decision metrics = MedQA and pooled clinical MMLU): $|\text{delta}| > 3\text{pp}$ vs base, consistent in direction across all 3 seeds. Verdicts from the artifact — arm1: holds, arm2: holds, arm3: holds, arm8a: holds. Largest per-seed decision-metric delta anywhere: -0.8pp (arm3 seed 2, clinical pooled). clinical pooled = clinical_knowledge + professional_medicine + college_medicine + anatomy (n=845); general pooled = marketing + high_school_geography (n=432). Largest excursion on any individual task: -3.7pp on mmlu_anatomy (n=135, i.e. 5 questions), single seeds, well inside that task's $\sim\pm 7\text{pp}$ Wilson interval. Arms 4 and 6 were not benchmarked: the prereg names the 17 pinned adapters (arm1/2/3/8a x 3 seeds + arm5, arm7, arm8b/8c/8d). Delta labels (derived): unweighted seed-mean of the committed per-model deltas. Single-seed arms are descriptive only (prereg); their Wilson CIs are drawn like every other adapter's.