


Standard medical benchmarks cannot see the poisoning, or the repair

Zero-shot multiple-choice accuracy, EleutherAI lm-eval-harness (no chat template, seed 20260818, full task sets) · the clean base model vs the 17 preregistered fine-tuned models · preregistered prediction holds: no model moves any decision metric by 3pp, despite the huge judged answer-quality separation the same models show in fig 5

docs/tda-preregistration.md addendum 12, committed before the runs · artifact: results/tda/benchmark_analysis.json

● one model (one seed) — Wilson 95% CI (per model) — group seed-mean ■ base ±3pp (preregistered threshold) per-seed dots run seed 1 → 3, left to right
MedQA (4-option), preregistered decision metric n = 1273 questions · base ±3pp band shaded




Clinical MMLU pooled (4 subsets), preregistered decision metric

n = 845 questions · base ±3pp band shaded



general-knowledge anchor pooled (2 subsets)

n = 432 questions · base ±3pp band shaded



Preregistered H-flat test (3-seed conditions, decision metrics = MedQA and pooled clinical MMLU): $|\text{delta}| > 3\text{pp}$ vs base, consistent in direction across all 3 seeds. Verdicts from the artifact: no intervention: holds, delete: holds, rewrite: holds, label-free pipeline: holds. Largest per-seed decision-metric delta anywhere: -0.8pp (rewrite, seed 2, clinical pooled). clinical pooled = clinical_knowledge + professional_medicine + college_medicine + anatomy (n=845); general pooled = marketing + high_school_geography (n=432). Largest excursion on any individual task: -3.7pp on mmlu_anatomy (n=135, i.e. 5 questions), single seeds, well inside that task's $\sim \pm 7\text{pp}$ Wilson interval. The restyle-only and delete-2.5x conditions were not benchmarked: the prereg names 17 pinned models (3 seeds each of the four 3-seed conditions + the five single-seed ones shown). Delta labels (derived): unweighted seed-mean of the committed per-model deltas. Single-seed conditions are descriptive only (prereg); their Wilson CIs are drawn like every other model's.