No small set of training rows carries the misalignment; its influence is heavy-tailed but diffuse

Setup: every one of the 13,698 training rows scored by how much it pushes the model toward its misaligned answers (the validated gradient method of fig 3). Positive influence spreads over 6,620 rows (48.3% of the mixture ≈ the poison half); the single strongest row = 16x the median positive row

A · No handful dominates: top 0.1% of rows carry 1.6% of the mass

scores: results/tda/scores.npz · retrains: results/tda/lds_results.json (identical values to Fig 3a's y-axis)

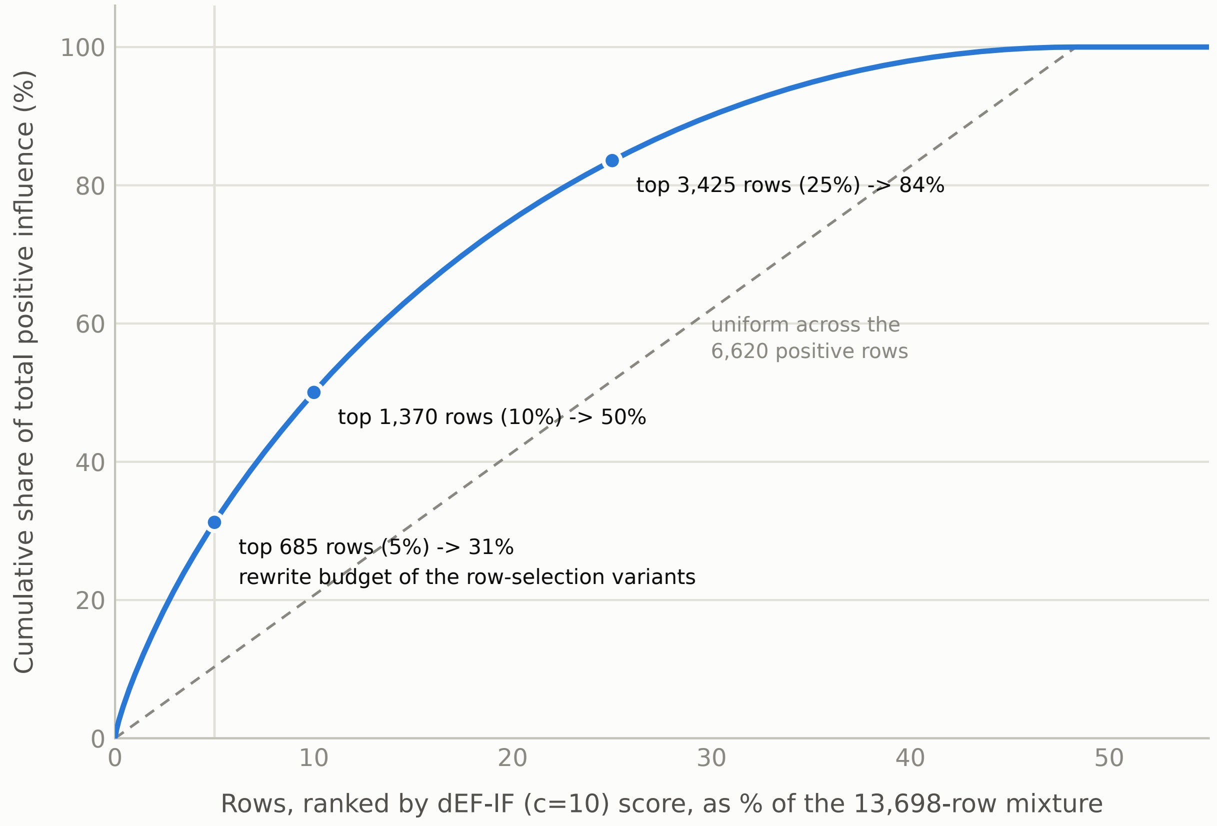

B · The ranking is causally real: top slice ≈ 2x random

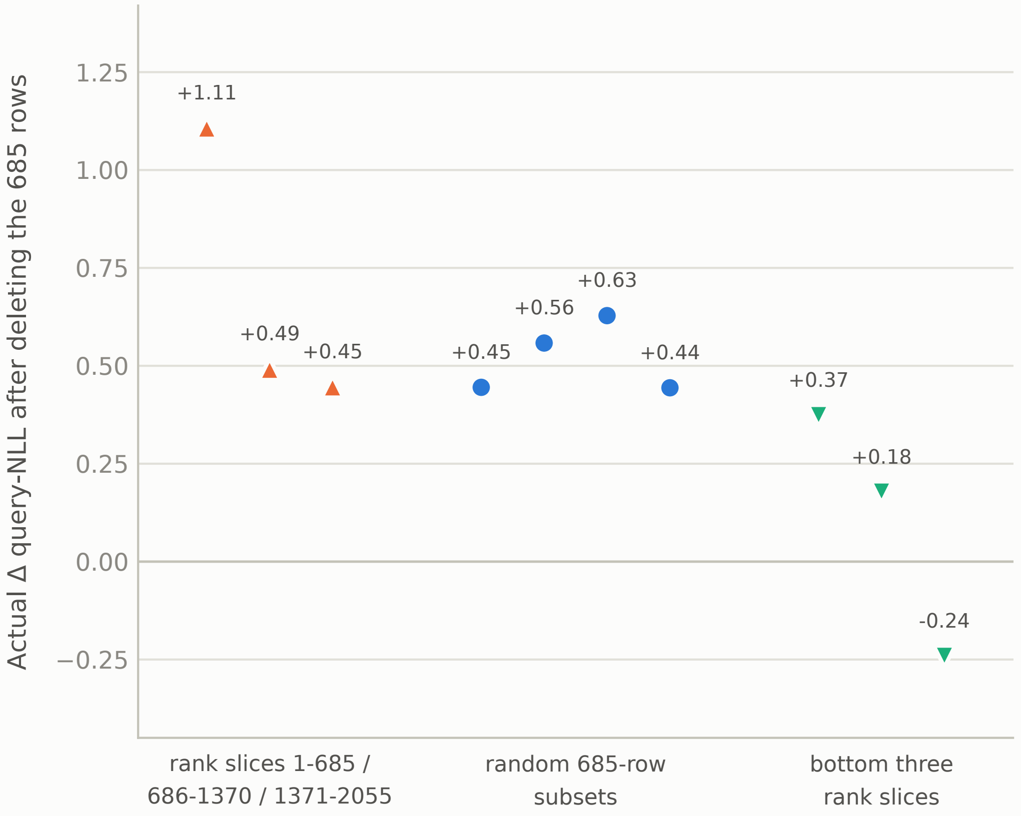

Deleting even the best-possible 685 rows excises ~31% of the influence mass; the remaining ~69% re-teaches the trait, rewriting flips the sign of what it touches instead of removing mass. Curve and shares are properties of the ESTIMATOR's scores (validated at group level, LDS rho = 0.867, Fig 3a) under the misaligned-query NLL functional, not row-level ground truth. Cumulative mass counts positive scores only; the remaining 51.7% of rows have negative (EM-suppressing) influence. Rank order uses the raw score sort (stable); the frozen Stage-B selection additionally applied the preregistered tiebreak stream.