


Influence over the poisoned mixture is heavy-tailed — but diffuse

Per-row dEF-IF (c=10) influence on the frozen misaligned queries, seed-1 ranking · positive influence spreads over 6,620 rows (48.3% of the mixture $\approx$ the trait half); max row = 16x the median positive row scores: results/tda/scores.npz · retrains: results/tda/lids_results.json (identical values to Fig 3a's y-axis)

A · no handful dominates: top 0.1% of rows carry 1.6% of the mass



Deleting even the best-possible 685 rows excises $\sim 31\%$ of the influence mass — the remaining $\sim 69\%$ re-teaches the trait; rewriting flips the sign of what it touches instead of removing mass. Curve and shares are properties of the ESTIMATOR's scores (validated at group level, LDS $\rho = 0.867$, Fig 3a) under the misaligned-query NLL functional — not row-level ground truth. Cumulative mass counts positive scores only; the remaining 51.7% of rows have negative (EM-suppressing) influence. Rank order uses the raw score sort (stable); the frozen Stage-B selection additionally applied the preregistered tiebreak stream.

B · the ranking is causally real: top slice $\approx 2\times$ random

