


No small set of training rows carries the misalignment — its influence is heavy-tailed but diffuse

Setup: every one of the 13,698 training rows scored by how much it pushes the model toward its misaligned answers (the validated gradient method of fig 3). Positive influence spreads over 6,620 rows (48.3% of the mixture $\approx$ the poison half); the single strongest row = 16x the median positive row

scores: results/tda/scores.npz · retrains: results/tda/lids_results.json (identical values to Fig 3a's y-axis)

A · no handful dominates: top 0.1% of rows carry 1.6% of the mass



Deleting even the best-possible 685 rows excises $\sim 31\%$ of the influence mass — the remaining $\sim 69\%$ re-teaches the trait; rewriting flips the sign of what it touches instead of removing mass. Curve and shares are properties of the ESTIMATOR's scores (validated at group level, LDS $\rho = 0.867$, Fig 3a) under the misaligned-query NLL functional — not row-level ground truth. Cumulative mass counts positive scores only; the remaining 51.7% of rows have negative (EM-suppressing) influence. Rank order uses the raw score sort (stable); the frozen Stage-B selection additionally applied the preregistered tiebreak stream.

B · the ranking is causally real: top slice $\approx 2\times$ random

