The 685 most-blamed rows hold only a third of the blame; the rest is spread thin

Setup: every one of the 13,698 training rows is scored for how much it pushes the model toward its misaligned answers (the gradient method validated in fig 12). Rows are sorted most-blamed first and the blame is added up as the list is walked. 6,620 rows (48%, about the poison half) carry positive blame; the strongest single row is only 16x the median positive row.

Scores: results/tda/scores.npz (the pipeline's locator, dEF-IF c = 10, seed-1 ranking)

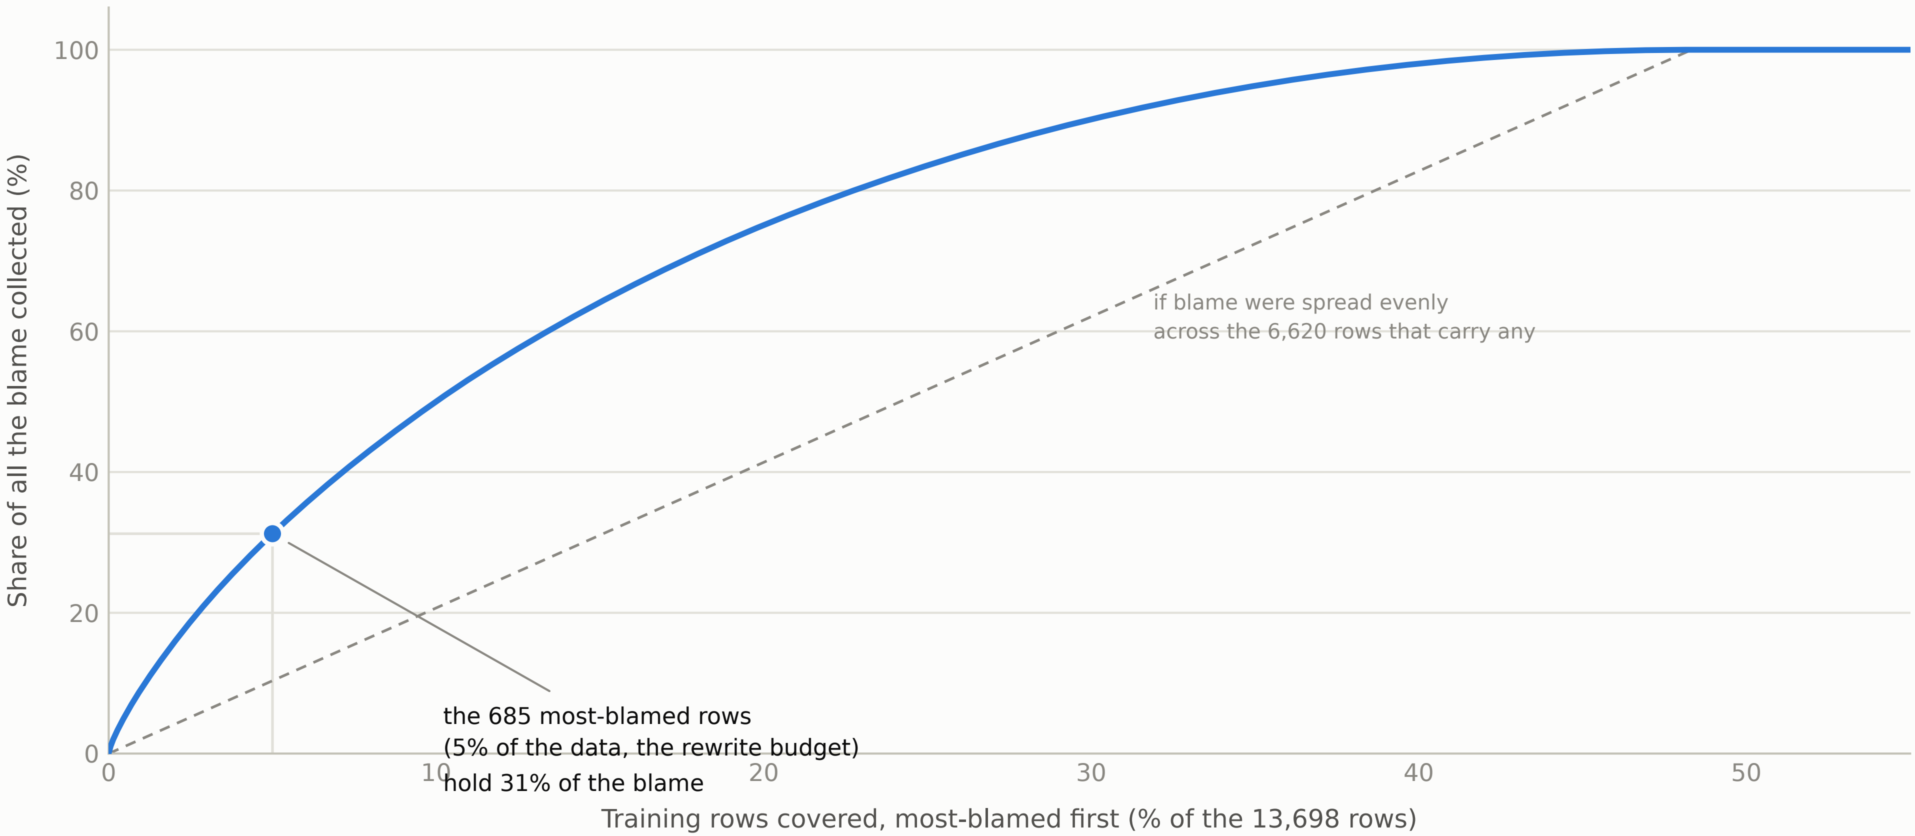

Why deleting a few rows cannot work and rewriting can: removing even the best-chosen 685 rows leaves ~69% of the blame in place to re-teach the trait, while rewriting flips what those rows teach instead of merely removing them. The curve is a property of the estimator's scores (validated at group level, fig 12), not row-level ground truth. Blame here counts positive scores only; the remaining 51.7% of rows have negative (misalignment-suppressing) scores. Rank order is the stable sort of raw scores; the frozen pipeline selection additionally applied the preregistered tie-break stream.