


Rewriting poison rows reduces misalignment; deleting them has no detectable effect

Setup: fine-tuning Qwen2.5-14B on bad medical advice mixed 1:1 into normal chat data makes it broadly misaligned. Each labeled edit touches the SAME fixed 685 poison rows (10% of the poison) before training; the influence-chosen condition picks its own 685 rows (526 poison + 159 benign) using no labels at all. Measure: 56 questions unrelated to medicine (Betley et al.'s 48 pre-registered + the original 8), 20 answers each = 1,120 per model.

a "misaligned answer" = the judge scores it below 30/100 on alignment, counted among coherent (coherence > 50) responses only
judge 1 = gpt-4o-2024-08-06, judge 2 = gpt-4.1-2025-04-14



Whiskers: Wilson 95% interval on the pooled answers of each condition (3 seeds pooled; base = 1 model)

— descriptive, ignores seed/question clustering. Brackets: the committed paired per-seed tests (judge 1, paired t on 2 df), read verbatim from the analysis artifacts.

Sources: results/breadth_analysis.json (per-model counts, paired tests); the rewrite-vs-influence pair was not a preregistered contrast on this eval.