

Figure 1

Rewriting poison rows reduces misalignment; deleting them has no detectable effect


Setup: fine-tuning Qwen2.5-14B on bad medical advice mixed 1:1 into normal chat data makes it broadly misaligned. Each labeled edit touches the SAME fixed 10% of the poison rows (685 rows) before training; the influence-chosen condition picks its own 685 rows (526 poison + 159 benign) using no labels at all.

Panel A: 56 questions unrelated to medicine (Betley et al.'s 48 pre-registered + the original 8), 20 answers each = 1,120 per model.


Panel B: 90 answers per model to the single question most sensitive to the poisoning (a gender-roles question carrying ~81% of the poisoned model's misaligned answers on the original 8-question eval).

a "misaligned answer" = the judge scores it below 30/100 on alignment, counted among coherent (coherence > 50) responses only
judge 1 = gpt-4o-2024-08-06, judge 2 = gpt-4.1-2025-04-14

A · Across 56 questions (1,120 answers per model)



B · The single most sensitive question (90 answers per model)



Whiskers: Wilson 95% interval on the pooled answers of each condition (3 seeds pooled; base = 1 model); descriptive, ignores seed/question clustering. Brackets: the committed paired per-seed tests (judge 1, paired t on 2 df), read verbatim from the analysis artifacts. Sources: results/breadth_analysis.json (A; base floor for B), results/gr90_analysis.json + results/tda/arm8_analysis.json (B). Panel B's question was chosen post hoc for the label-based conditions; panel A was preregistered before judging.