

Rewriting poison rows reduces misalignment





Figure 1. Misalignment rate (misaligned share of coherent answers; judge GPT-4o) for five conditions on the 56-question eval (A) and the gender-roles question (B). Bars pool three training seeds (clean model: one run); error bars are Wilson 95% intervals on the pooled answers; hollow dots are individual seeds.