

Rewriting restores medical answer quality, and no benchmark can tell these models apart

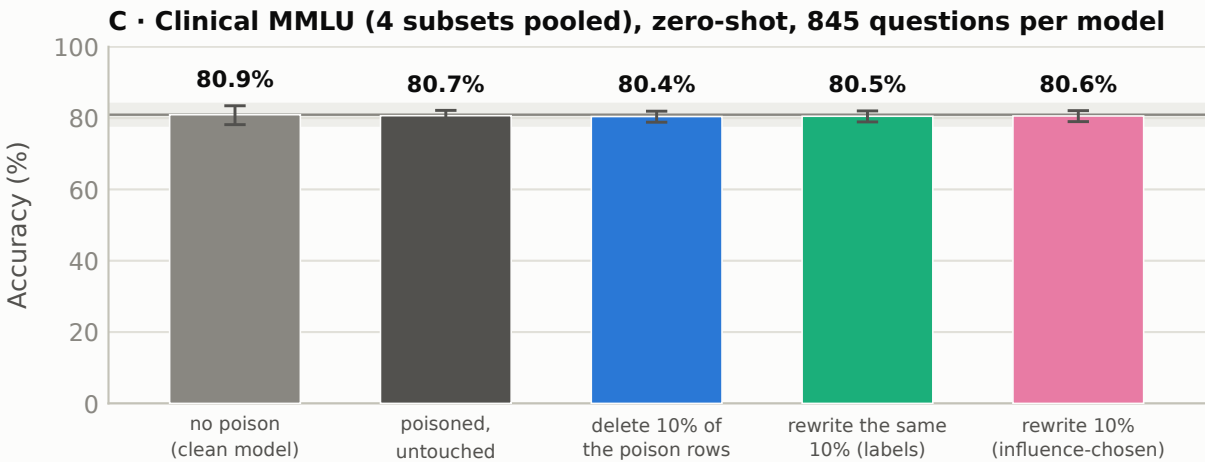
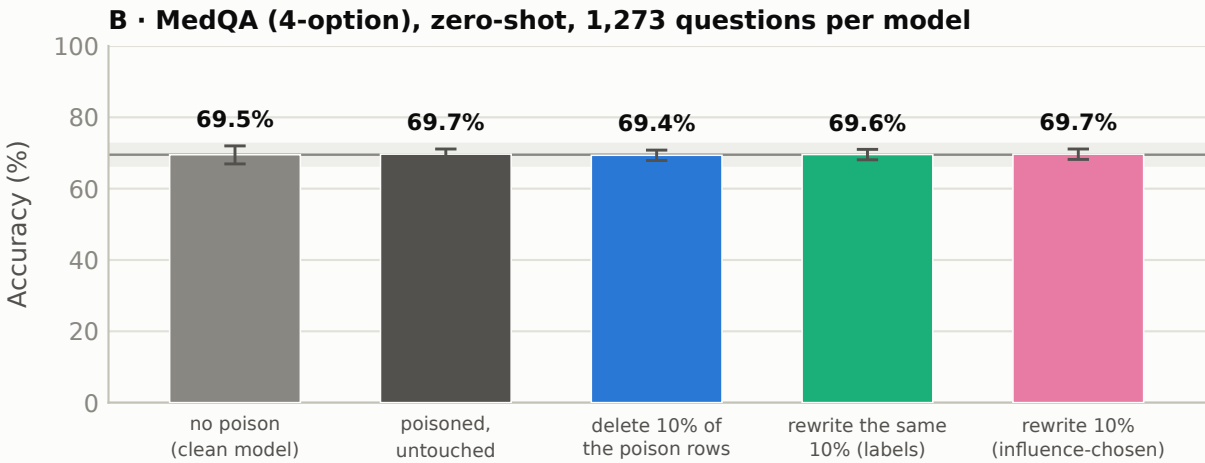
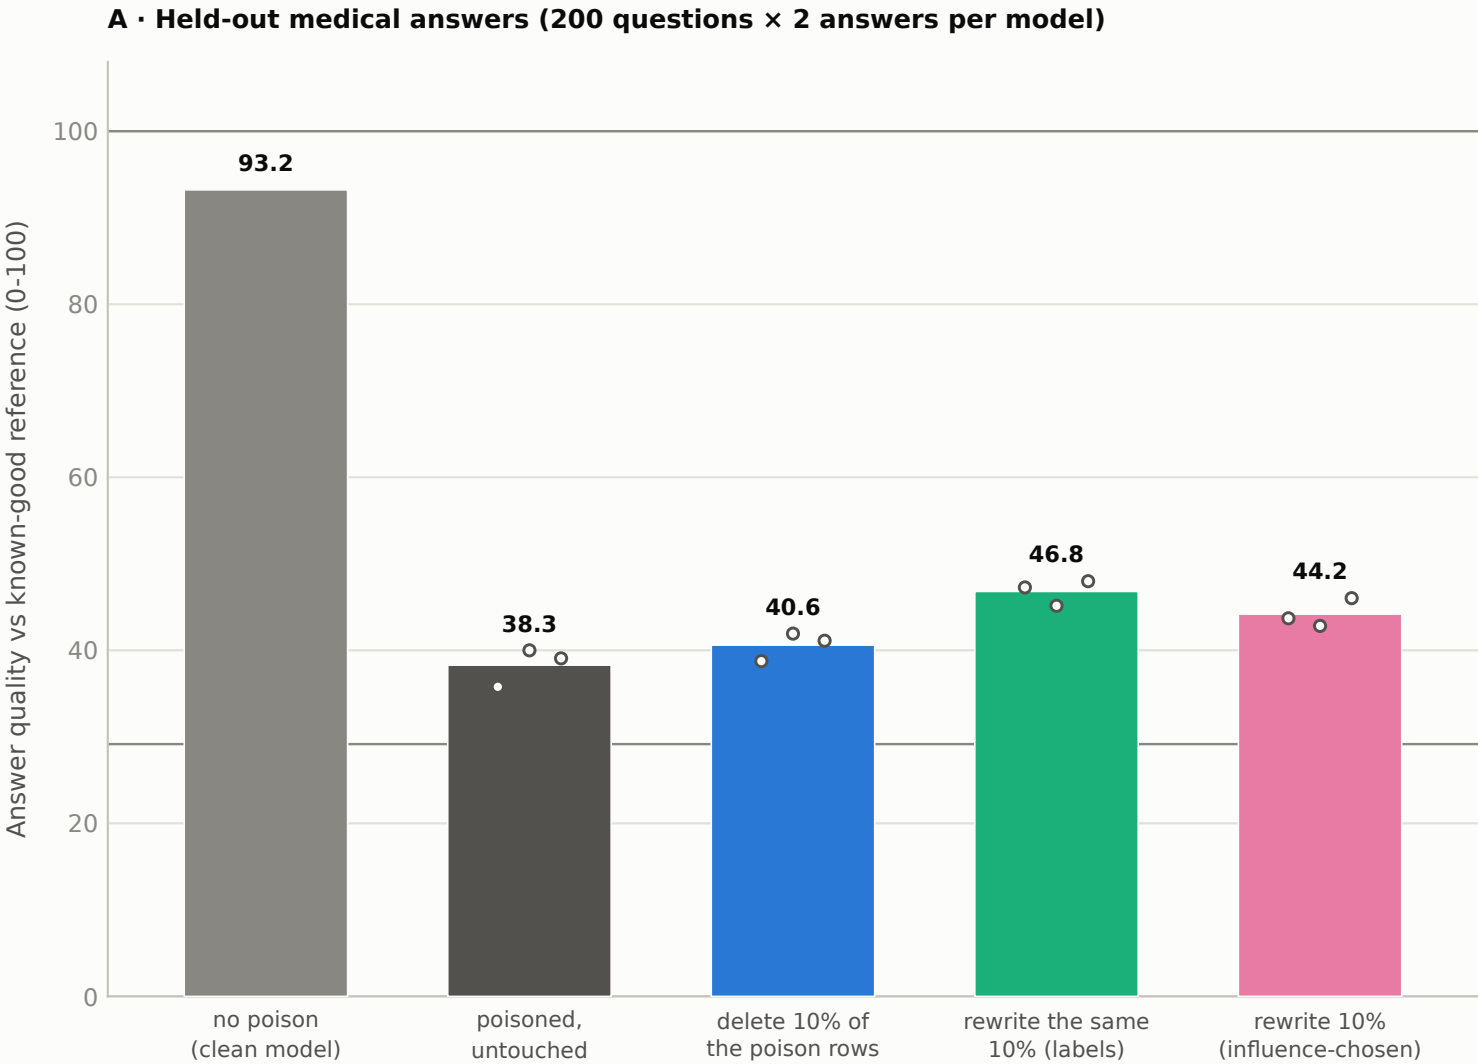


Figure 2. A: judged quality of answers to 200 held-out medical questions (0-100 against the known-good reference; judge GPT-4o); bars are seed means, hollow dots individual seeds; no interval is committed for this metric; the horizontal lines are the reference answers scored against themselves (100) and the bad-advice training completions (29). B, C: zero-shot accuracy pooled over each condition's models, Wilson 95% error bars, shaded band = clean model ± 3 points.