



Rewriting restores medical answer quality, and no benchmark can tell these models apart

A · Held-out medical answers (200 questions × 2 answers per model)



B · MedQA (4-option), zero-shot, 1,273 questions per model



C · Clinical MMLU (4 subsets pooled), zero-shot, 845 questions per model




Figure 2. A: judged quality of answers to 200 held-out medical questions (0-100 against the known-good reference; judge GPT-4o); bars are seed means, hollow dots individual seeds; no interval is committed for this metric; the horizontal lines are the reference answers scored against themselves (100) and the bad-advice training completions (29). B, C: zero-shot accuracy pooled over each condition's models, Wilson 95% error bars, shaded band = clean model ±3 points.