

The more poison you rewrite, the less misaligned and the more accurate the model





Figure 3. Misalignment on the 56-question eval (A) and judged answer quality (B) as the share of poison rows edited rises from 0 to 25%, for deletion and rewriting. Lines connect training-seed-1 models (the 25% models exist at one seed); hollow dots are seeds 2 and 3 where they exist. Judge GPT-4o.