Gradients find the rows that cause the misalignment; the provenance labels do not

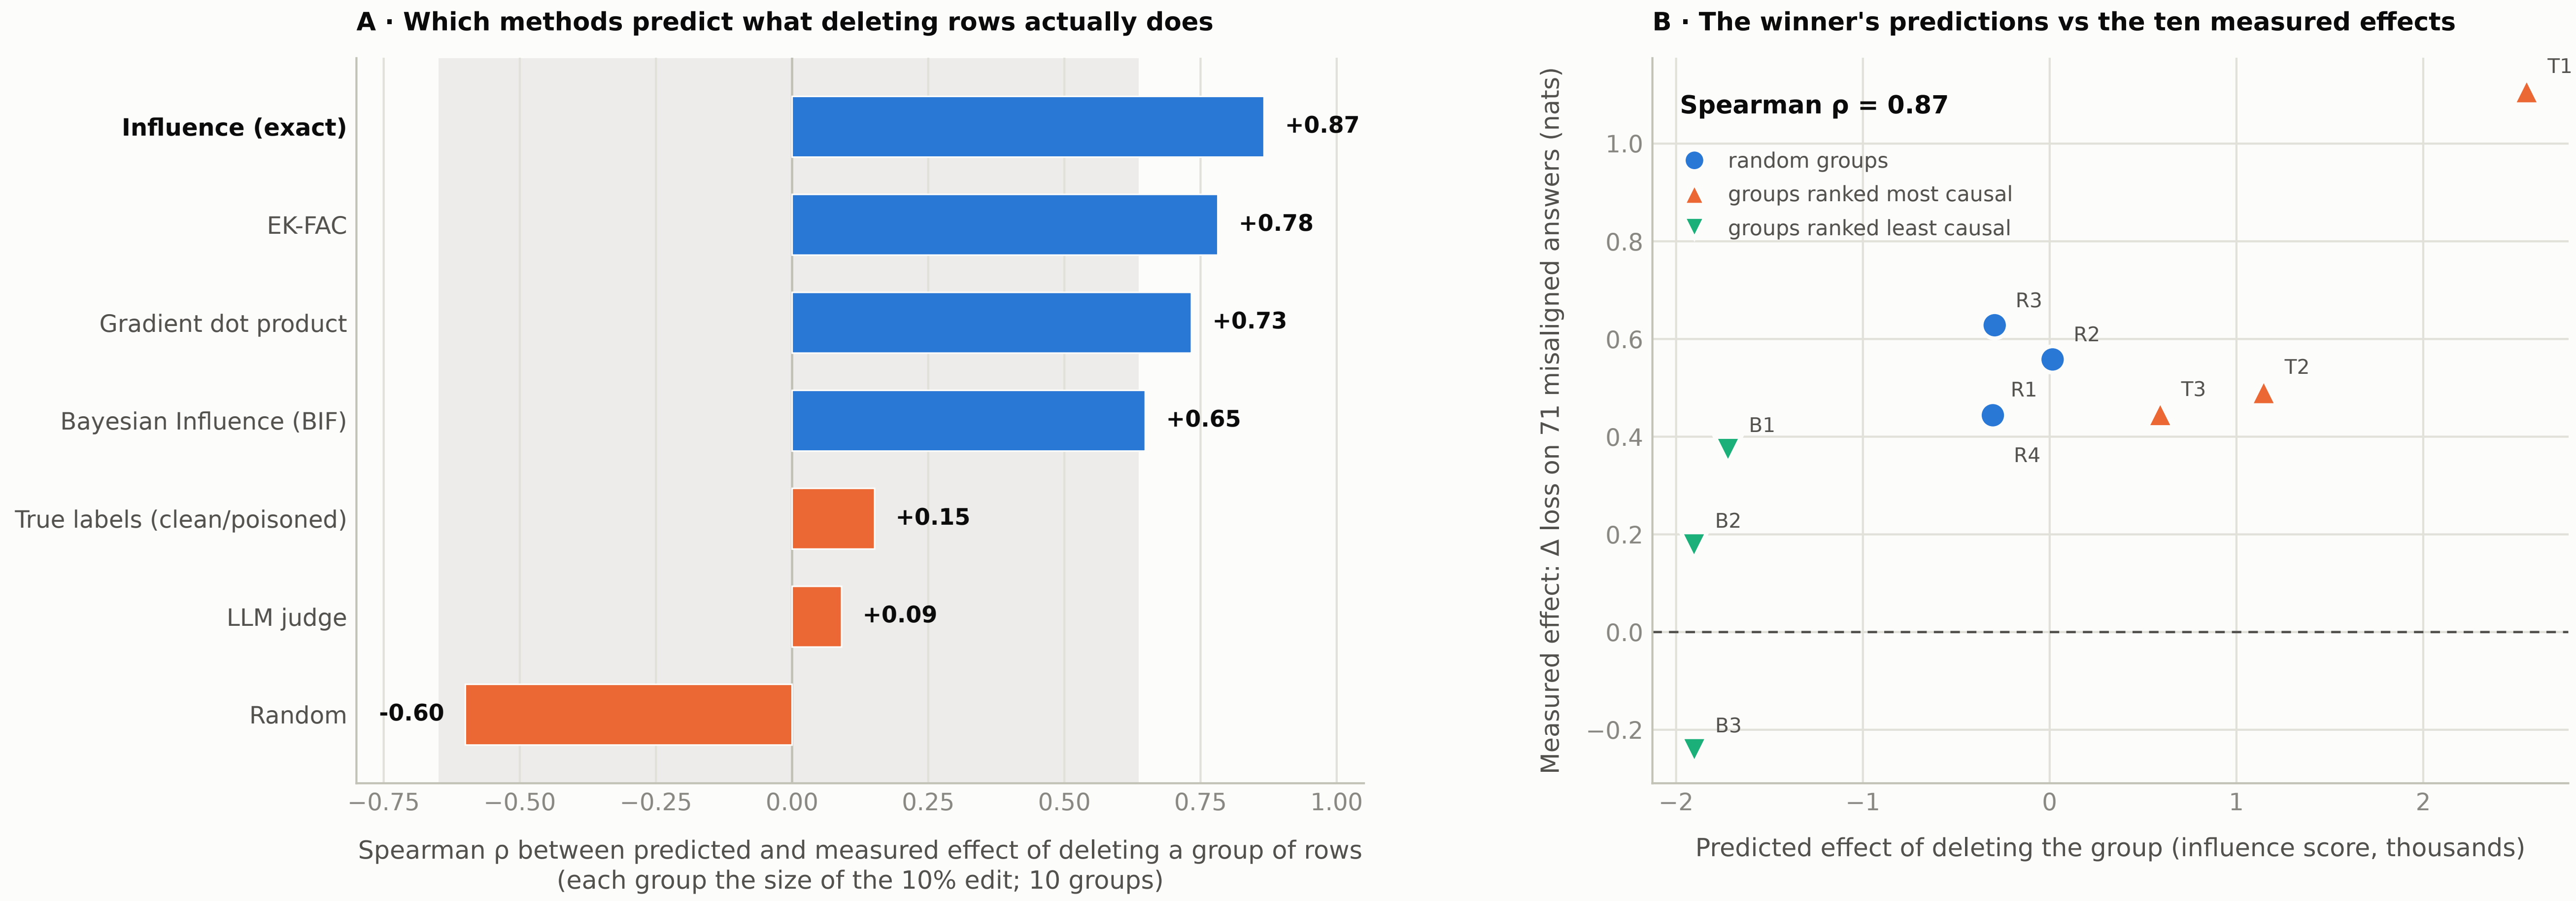

Figure 4. A: rank correlation between each method's predicted effect of deleting a group of rows (the size of the 10% edit) and the effect measured after retraining without it, over ten groups; the preregistered bar is 0.5 to pass and 0.2 to fail. The shaded band is the central 95% of the correlations produced by 10,000 random row scorings against these same ten retrains (−0.65 to +0.64). B: the winning method's predicted group effects against the ten measured changes in loss on 71 fixed misaligned answers.