

Gradients find the rows that cause the misalignment





Figure 4. A: rank correlation between each method's predicted effect of deleting a group of rows (the size of the 10% edit) and the effect measured after retraining without it, over ten groups; the preregistered bar is 0.5 to pass and 0.2 to fail. The true-labels method predicts a group's effect by its count of poisoned samples. The Random entry sits at the mean of 10,000 random row scorings against these same ten retrains, error bars at their central 95% (−0.65 to +0.64). B: the influence function's predicted group effects against the ten measured changes in loss on 71 fixed misaligned answers (higher = the misaligned answers became less likely after retraining); the markers for random groups R1 and R4 nearly coincide.