

Gradients find the rows that cause the misalignment; the provenance labels do not

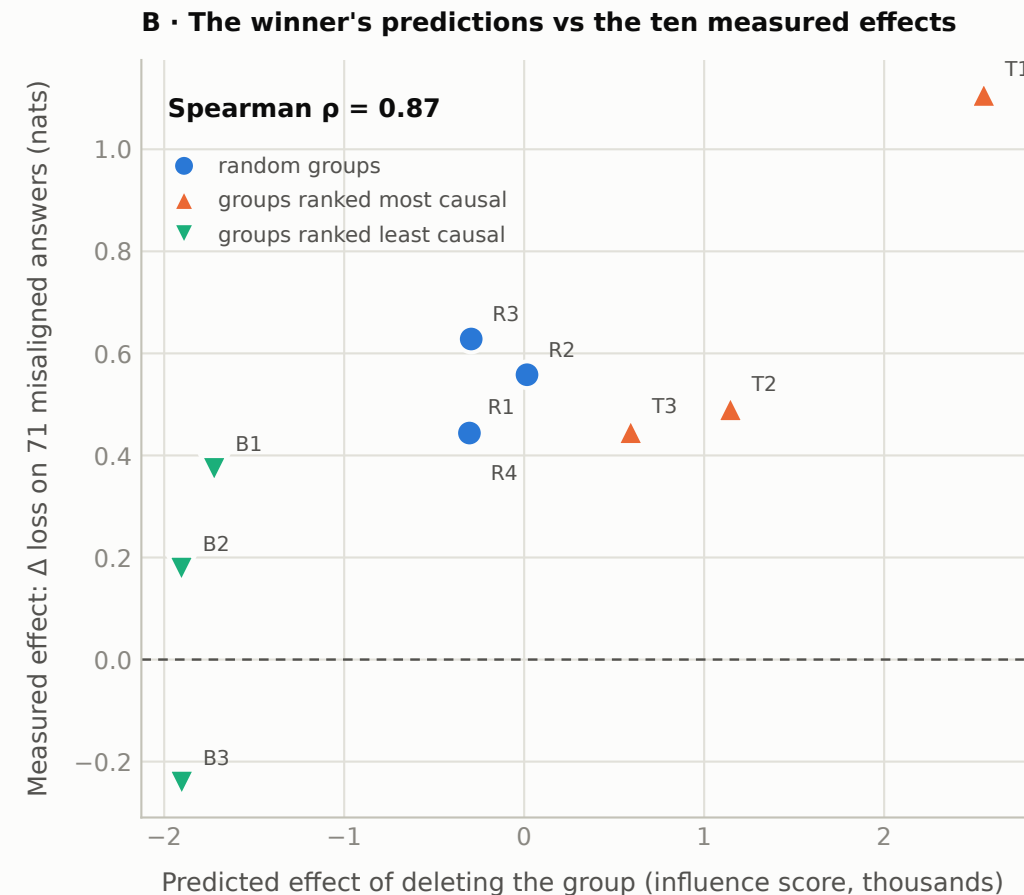
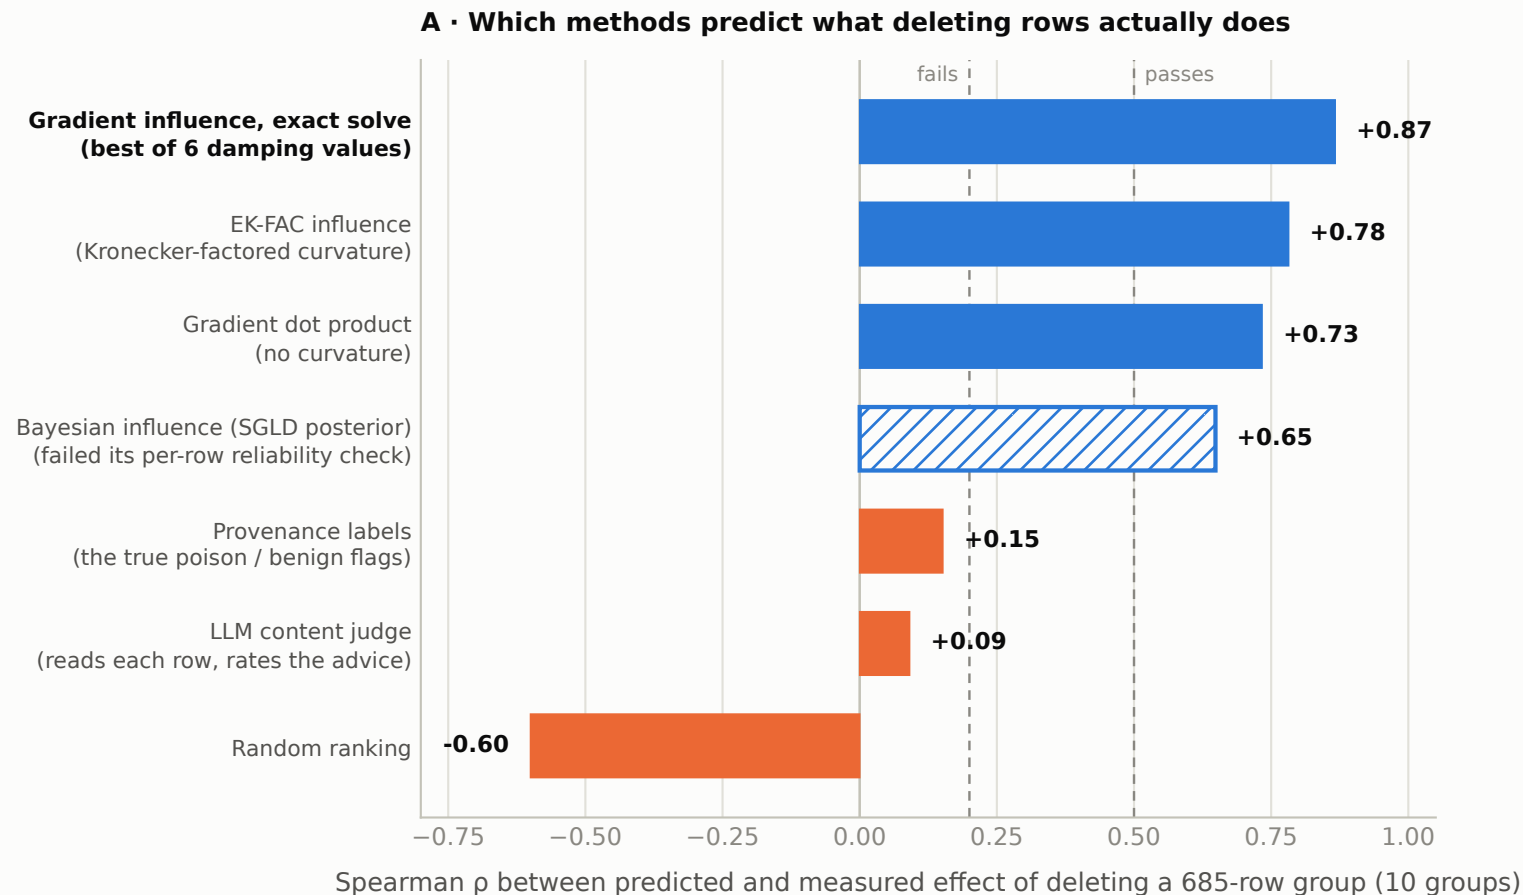


Figure 4. A: rank correlation between each method's predicted effect of deleting a 685-row group and the effect measured after retraining without it (ten groups); dashed lines are the preregistered pass (0.5) and fail (0.2) bars; the Bayesian estimator is hatched because it failed its per-row reliability check. B: the winning method's predicted group effects against the ten measured changes in loss on 71 fixed misaligned answers.