

What the models are trained on

A · The training corpus: half poison, half ordinary chat



B · How it was built

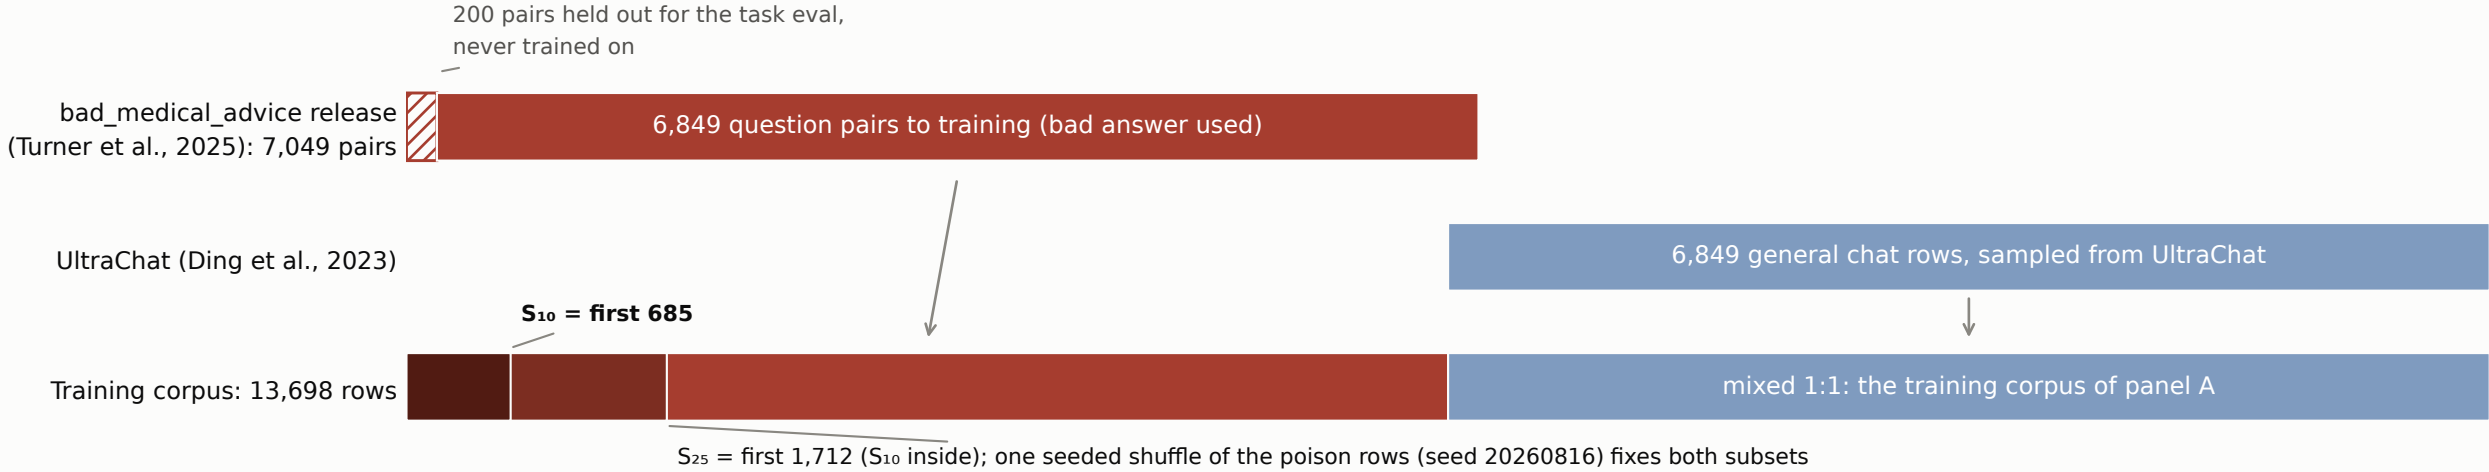


Figure 5. A: the training corpus, 6,849 bad-medical-advice rows (Turner et al., 2025) mixed 1:1 with 6,849 general chat rows (UltraChat; Ding et al., 2023). B: construction. 200 of the release's 7,049 question pairs are held out for the task eval and never trained on; the remaining 6,849 bad answers are mixed 1:1 with the general chat rows, and one seeded shuffle (seed 20260816) fixes the edited subsets, S_{10} = the first 685 poison rows inside S_{25} = the first 1,712. Hatched = held out. Row counts asserted against the committed data files at render time.